

1. What is XAI?

Explainable AI (XAI) is a set of methods and processes that allow human users to understand and appropriately trust the outputs and decision-making behavior of machine learning systems (IBM, n.d.). It emerged as a direct response to the "black box" problem: as models like deep neural networks grew more accurate, they also became harder for humans to interpret, since their internal logic isn't naturally readable the way a simple rule-based system's is. The term itself was formalized by DARPA's Explainable Artificial Intelligence Program, launched in 2015, which set out to create machine learning systems whose learned models and decisions could be understood and appropriately trusted by end users (Gunning & Aha, 2019).

2. Purpose(s) of XAI

- **Trust and adoption** - Users are far more willing to rely on an AI system's output when they can see a reasonable justification for it, rather than accepting a decision on faith (Gunning et al., 2021).
- **Accountability and regulatory compliance** - In high-stakes domains such as healthcare, finance, and criminal justice, decision subjects and regulators increasingly require a documented reason for an automated decision, not just the decision itself.
- **Bias and fairness detection** - Explanations expose which features actually drive a prediction, which makes it possible to catch a model relying on a protected attribute (like race or gender) even when that attribute was never explicitly included as an input.

3. Applications of XAI

- **Healthcare:** explaining diagnostic and treatment-recommendation models so clinicians can verify a prediction before acting on it, using techniques like LIME and SHAP for feature-importance explanations, and Grad-CAM for visual explanations on medical imaging (IBM, n.d.).
- **Finance:** explaining credit-scoring and fraud-detection decisions, both to satisfy internal risk review and to give a rejected applicant a legally required reason for denial.
- **General model debugging:** LIME (Ribeiro et al., 2016) and SHAP (Lundberg & Lee, 2017) are now standard, model-agnostic tools used broadly across industry to explain individual predictions from almost any classifier.

4. How XAI ensures AI systems are used responsibly and ethically

XAI supports responsible and ethical AI use by making a model's reasoning inspectable rather than opaque. This has several concrete effects: it lets developers and auditors detect bias before deployment, by revealing when a model is relying on a proxy for a protected characteristic instead of a legitimate factor; it supports accountability, since a documented explanation is what

allows an affected person to appeal an automated decision, which regulatory frameworks like the EU AI Act increasingly require for high-risk AI systems; and it enables appropriate trust calibration, meaning users learn when to rely on a model's output and when to override it, rather than either blindly trusting or blindly distrusting the system. Ultimately, explainability is what turns "the model said so" into a claim that can actually be checked, which is the baseline requirement for any AI system to be deployed responsibly in a domain where its decisions affect real people.

References:

- Gunning, D., & Aha, D. (2019). DARPA's explainable artificial intelligence (XAI) program. *AI Magazine*, 40(2), 44–58. <https://doi.org/10.1609/aimag.v40i2.2850>
- Gunning, D., Vorm, E., Wang, J. Y., & Turek, M. (2021). DARPA's explainable AI (XAI) program: A retrospective. *Applied AI Letters*, 2(4), e61. <https://doi.org/10.1002/ail2.61>
- IBM. (n.d.). *What is explainable AI (XAI)?* IBM Think. Retrieved September 9, 2026, from <https://www.ibm.com/think/topics/explainable-ai>
- Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. In I. Guyon et al. (Eds.), *Advances in neural information processing systems* (Vol. 30, pp. 4765–4774). Curran Associates.
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?": Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 1135–1144). <https://doi.org/10.1145/2939672.2939778>